

Check-IA – Evaluation Methodology Note

Version 1 · July 2026

PURPOSE

This note describes the method and results of the first systematic evaluation of Check-IA, the information-verification platform for the Sahel region. The evaluation covers verification of French-language text, measured on the system as deployed in production.

CORPUS

396 real claims were built from verifications published between 2019 and 2026 by two independent Malian fact-checking organizations: **Benbere** (#BenbereVerif) and **Voix de Mopti** (Mopti Check). Each claim was rephrased into the neutral form a user would submit (with no verdict vocabulary), with human spot-checking of the rephrasings' fidelity. 371 claims have a strictly **true-or-false** reference verdict; the remaining 25 belong to nuanced categories (**misleading, out of context**) and are measured separately.

GROUND TRUTH

Each claim's reference verdict is anchored in the conclusion published by the original fact-checker. Because reality evolves — some claims that were false when verified later came true — the entire corpus underwent a temporal audit: every claim whose status could plausibly have changed was re-verified against events through July 2026, with sources, and reference verdicts were updated accordingly. The fact-checkers named above are the authors of the original verifications; they did not supervise or endorse this evaluation.

PROTOCOL

Each claim was submitted to the production verification system under the same conditions as a user submission. The verdicts returned were compared against the reference verdicts. A subset was also submitted end-to-end through the deployed application to measure real response times.

PUBLISHED METRICS

- **Correct verdict / caution / wrong verdict:** the distribution of outcomes across the 371 true/false claims. "Caution" designates cases where the system chooses not to rule rather than guess.
- **Accuracy of decisive verdicts:** the share of correct verdicts among claims where the system did rule.
- **False claims wrongly validated:** the share of false claims to which the system assigned a "true" verdict — the most serious error a verification tool can make.
- **Median response time:** measured end-to-end on the deployed application.

RESULTS (V1 – JULY 2026)

Measure	Result
Correct verdict	72.8% (270/371)
Caution (no verdict)	24.3% (90/371)
Wrong verdict	3.0% (11/371)
Accuracy of decisive verdicts	96.1% (270/281)
False claims wrongly validated	1.7% (6/355)
Median response time	≈ 30 seconds

On the 25 nuanced claims (misleading, out of context), the system concludes "false" in 60% of cases — an acceptable approximation for a three-verdict system, measured separately so as not to distort the headline figures.

LIMITATIONS

- The corpus is mostly composed of false claims (355 of 371), reflecting the fact-checking publications it is drawn from. Measures relating to true claims rest on a small sample and are indicative.
- Claims about the identity of images or videos ("this photo shows...") have limitations of their own when evaluated through the text pathway.
- Results reflect the state of the system and of the web as of July 2026.

VERSIONS

Evaluations are dated and versioned. Results of subsequent evaluations will be published as v2, v3, ... on the Reliability page at check-ia.app.